

LLM Selection Strategy for Real Estate AI Platform

Objective

The Real Estate AI platform supports multiple AI workloads across both the ingestion pipeline and retrieval pipeline. These workloads have different latency, reasoning, accuracy, and cost requirements.

Instead of selecting a single LLM for the entire platform, I would design a multi-model architecture with intelligent routing.

The guiding principle is:

Use the cheapest model that can reliably perform the task.

Where LLMs Are Used

Ingestion Pipeline

Document Classification

Examples:

- Sale Deed
- Title Deed
- Property Brochure
- CRM Notes
- Email
- Legal Contract

Requirements:

- Low latency
- Structured classification
- High throughput

Recommended Model Tier:

- Small LLM
 - Fine-tuned classifier
-

Metadata Extraction

Examples:

- Property Type
- Price
- Location
- Registration Number
- Ownership Information

Requirements:

- Structured output
- Function calling
- Low cost

Recommended Model Tier:

- Small LLM
-

OCR Enhancement and Vision Understanding

Examples:

- Floor Plans
- Property Images
- Site Maps
- Scanned Legal Documents

Requirements:

- Multimodal support
- Image understanding

Recommended Model Tier:

- Vision Language Model (VLM)
-

Section Summarization (Optional)

Examples:

- Section Summary
- Parent Node Summary

Requirements:

- Offline processing
- Cost optimization

Recommended Model Tier:

- Small or Medium LLM
-

Query Pipeline

Query Understanding

Examples:

- Intent Detection
- Entity Extraction
- Query Classification
- Metadata Extraction

Requirements:

- Low latency
- High throughput

Recommended Model Tier:

- Small LLM
-

Query Planner

Responsibilities:

- SQL Routing
- Vector Search Routing
- Hybrid Search Routing
- Legal QA Routing

Requirements:

- Lightweight reasoning
- Structured output

Recommended Model Tier:

- Small LLM
-

Context Guardrails

Responsibilities:

- Prompt Injection Detection
- PII Detection
- Compliance Validation
- Tenant Isolation Validation

Requirements:

- Low latency
- High accuracy

Recommended Model Tier:

- Small LLM
-

Response Generation

Examples:

- Property Search Q&A
- Property Comparison
- Legal Question Answering
- Contract Interpretation
- Multi-document Analysis

Requirements:

- Strong reasoning
- Long context handling
- Citation generation
- Low hallucination rate

Recommended Model Tier:

- Medium or Large LLM
-

LLM Selection Criteria

1. Multilingual Support

The platform serves:

- English
- Hindi
- Kannada
- Tamil
- Telugu

The model must support:

- Multilingual reasoning
- Cross-lingual understanding

Example:

Query in Kannada

Documents in English

The model should still answer correctly.

2. Long Context Support

Documents may include:

- Sale Deeds
- Property Contracts
- Legal Agreements
- Property History

The model must support large context windows.

3. Tool Calling

The model should support:

- SQL Generation
- API Calls
- Query Planning

- Structured Output
-

4. Cost Efficiency

Metrics:

- Cost per query
 - Cost per tenant
 - Cost per feature
-

5. Latency

Target:

- Interactive search: < 2 seconds
 - Legal workflows: < 5 seconds
-

6. Hallucination Rate

The platform should prioritize:

Accuracy > Speed

Especially for legal and compliance workflows.

Recommended LLM Stack

Small Model Tier

Use Cases:

- Query Classification
- Metadata Extraction
- Query Planner
- Guardrails
- SQL Generation

Recommended Models:

Gemini 2.5 Flash

Pros:

- Fast
- Low cost
- Excellent multilingual support
- Strong tool calling

GPT-5 Mini

Pros:

- Strong structured output
- Good reasoning-to-cost ratio

Llama 4 Scout

Pros:

- Self-hosted option
- Lower infrastructure cost at scale

Recommended Choice:

Gemini 2.5 Flash

Medium Model Tier

Use Cases:

- Property Recommendations
- Summarization
- Context Compression
- Multi-step Search Workflows

Recommended Models:

Gemini 2.5 Pro

Pros:

- Strong multilingual capabilities
- Large context window
- Cost-effective

Claude Sonnet

Pros:

- Excellent document understanding
- Strong summarization

Recommended Choice:

Gemini 2.5 Pro

Large Model Tier

Use Cases:

- Legal QA
- Contract Interpretation
- Risk Analysis
- Multi-document Reasoning
- Complex Property Comparison

Recommended Models:

GPT-5

Pros:

- Excellent reasoning
- Strong agent workflows
- Superior tool use

Claude Opus

Pros:

- Exceptional legal document reasoning
- Long-form analysis

Gemini 2.5 Pro

Pros:

- Large context support
- Strong multimodal capabilities

Recommended Choice:

GPT-5 or Claude Opus based on benchmark results.

Model Routing Strategy

Instead of a single model:

User Query ↓ LLM Router ↓

Simple Tasks → Gemini 2.5 Flash

Medium Tasks → Gemini 2.5 Pro

Complex Legal Tasks → GPT-5 / Claude Opus

Examples:

Route to Small Model

- Metadata Extraction
- SQL Generation
- Query Planning
- Intent Classification

Route to Medium Model

- Summarization
- Property Recommendation
- Search Assistance

Route to Large Model

- Legal Reasoning
 - Contract Analysis
 - Multi-document QA
-

Evaluation Before Final Selection

I would build a benchmark dataset containing:

Property Search

Example:

"Show east-facing apartments under 1 crore."

Legal Questions

Example:

"Can ownership be transferred to a family member?"

Cross-Lingual Queries

Example:

Kannada query with English documents.

Multi-document Reasoning

Example:

Compare these three properties.

Metrics:

- Accuracy
- Faithfulness
- Hallucination Rate
- Citation Accuracy
- Latency
- Cost per Query
- User Satisfaction

Final Recommendation

I would not standardize on a single LLM.

Instead:

- Gemini 2.5 Flash for lightweight tasks
- Gemini 2.5 Pro for medium-complexity workflows
- GPT-5 or Claude Opus for legal reasoning and complex multi-document analysis

This architecture balances:

- Accuracy

- Latency
- Cost
- Scalability
- Enterprise Reliability

while providing the flexibility to evolve as model capabilities improve.